

Group project in TDT4215

Web-intelligence

Alessio DE ANGELIS Martin HENGSTBERGER Simon STASTNY

April 20, 2013

Contents

1	Introduction	2
2	System architecture	2
3	Components function and role	2
3.1	Parsing ICD and ATC files	2
3.2	Tools for indexing and querying	3
3.3	Parsing patient cases	4
3.4	Parsing the gold standard file	4
3.5	Parsing the book	4
4	Limitations given in the assignment. Present theories used. Are any theories outside of the curriculum been used? Explain the reason for the selection	5
5	How have the theories been adapted to the assignment?	5
6	Summary of the results.	5
7	Evaluation and discussion of the results. Limitations of your evaluation results.	5
8	Thoughts of potential improvements.	5
9	Conclusions.	5
10	Future work	5

1 Introduction

FIXME

2 System architecture

Since the nature of most of the tasks in this software system is highly time consuming, we decided to take advantage of provided multithreading facilities that Java offers and we implemented the software using Java Executor Framework. This choice made the code a bit more complicated than a code of regular single-threaded application, but we managed to reduce the execution time significantly on multi-processor machine.¹

Given the fact that our main architectural drivers were short time to market and developer inexperience, we decided to use just basic facilities of chosen frameworks and tools. Spending too much time learning details of tools such as *Jena* and *Lucene* would leave with a little time to make the actual implementation.

As stated above, we are using some tools such as *Jena* and *Lucene*. Their particular use in the application is described below in corresponding chapters. Part from those tools, we are for example using SAX and Apache Tika for parsing tasks.

3 Components function and role

Many tasks occurring in the application share some similarities in behaviour and structure and thus we made the system very modular and easily modifiable. The sequence diagram shown in 3 displays a naïve view over the system. Please note, that **Indexer**, **Parser** and **QueryEngine** are general names for different implementations (ATC files are parsed and indexed differently than ICD files, however the structure of the software component is similar).

3.1 Parsing ICD and ATC files

There are several ways to parse ICD and ATC files to extract information from them: you can consider those as normal XML files, use regular expressions and so on.

We have chosen a semantic web tool, *Jena*.² As we can read from the website: *Jena is an API for reading, processing and writing RDF data in XML, N-triples and Turtle formats; an ontology API for handling OWL and RDFS ontologies; a rule-based inference engine for reasoning with RDF and OWL data sources; stores to allow large numbers of RDF triples to be efficiently stored on disk.*

This tool may seem to be a little bit verbose to use, but let us handling in a very efficient way both the ontologies, the ICD that is an OWL format and the ATC that is a turtle, without the risk of losing information from them. Moreover, we can have easily access to the information about all the rdf graph: we can keep track for example of the hierarchies, the subclasses, the synonyms and so on. All this extra information may be used to improve our query engine, doing some refinements when the pure textual search is not accurate enough.

When we parse ICD we store in java beans the values of the following properties **code_compacted** (the code of the disease), **label** (description of the disease), **underterm** (extra information about the disease, not every instance has it), **synonyms** and **subClassOf**. When we parse ATC we store in java beans the values of the code, the **label** and **subClassOf**.

¹Execution time on a machine with Intel Core i3 (M350 2.27GHz) was reduced by a factor of cca 3.

²<http://jena.apache.org/>

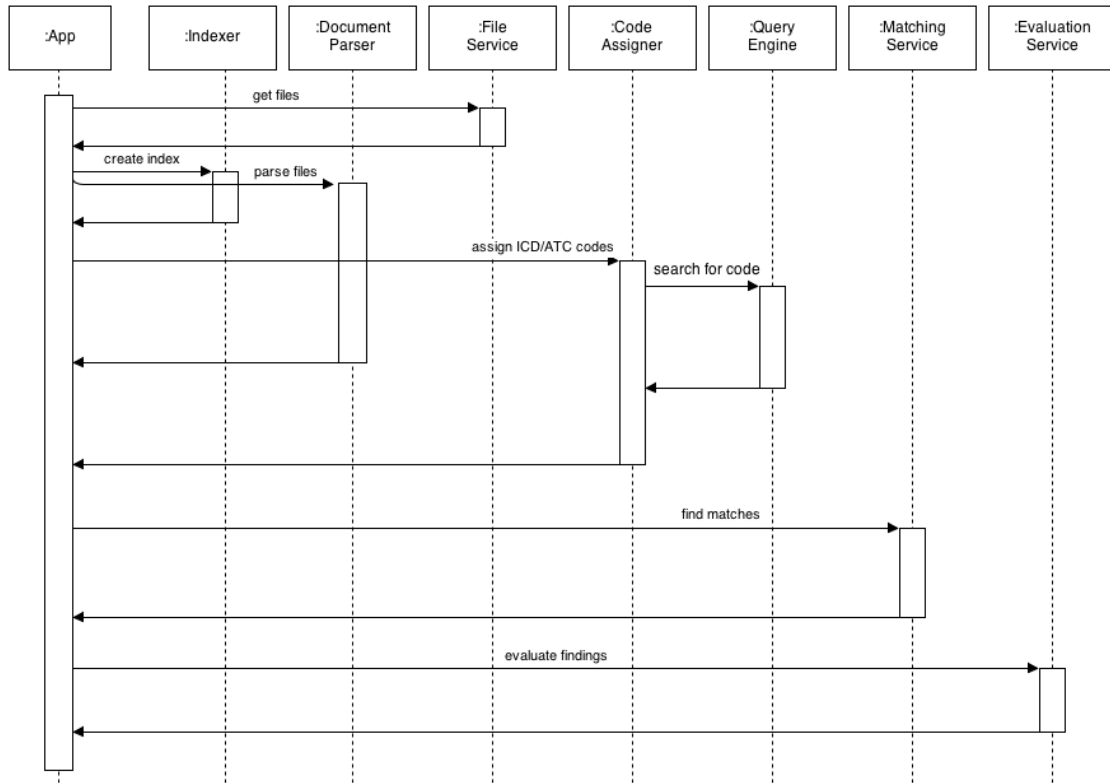


Figure 1: Sequence diagram




Figure 2: Class diagram

3.2 Tools for indexing and querying

For indexing and querying we are using an information retrieval approach, through a bag of words model. We are using *Lucene*,³ (TODO: if you want you can describe a little bit about *Lucene*, reading from the website). At the beginning we were trying to create an our own information retrieval engine, but we came across several problems. First of all how to deal with Norwegian language, our background is very poor since we are all international students in the group, how to stem correctly words or divide properly the *aggregates* (four, five words merged all together in one). *Lucene* is a library made by experts and their implementation of the *NorwegianAnalyzer* (with the Porter stemming algorithm), even if maybe not as accurate as the English one, has been revealed very useful for our purposes. Moreover it natively supports the vector space model for a more accurate retrieval than boolean.

³<http://lucene.apache.org/core/>

We keep two different indexes, one for the ATC documents and the second for the ICD documents. Since the indexing is very fast, we keep both the indexes in ram, for performance reasons during the queries (you don't need to access the disk that is usually time consuming). The ATC documents are composed by two fields: one for the `atc_code`, that just indexed without any further string processing; the second field, `label`, contains the description of the code and is the one used for extracting the documents. This field is indexed, stored, tokenized and calculates the term vector matrix.

The ICD documents are indexed according to three fields. The first one is for the property `icd:code_compacted`, to identify uniquely the codes. The second one is for the `label`: the field has the same properties as the ATC label. The third field contains all the terms of the label plus the other information for doing query expansion, such as the synonyms and the property `icd:underterm`. Adding those extra words let have a more accurate result when we submit the query to our engine. In fact at the moment we query the ICD index looking into the terms contained in the extra field.

3.3 Parsing patient cases

Since the assignment description states that preprocessing of patient case input file is legal, we took te chance and manually made an XML file representing the data contained in the `.doc` file provided. Parsing a `.doc` file would be a waste of time with no learning outcome. This as well gave us the opportunity to decide how the text is divided into sentences manually.

To parse this created XML document, we were using SAX parser with own implementation of `DefaultHandler` to produce `MedDocument` instances from the files read. This part of the work gets done in `PatientCaseParser`.

3.4 Parsing the gold standard file

To perform the evaluation easily, we decided to make an XML file for the gold standard file⁴ and let the application read it and compare the two results set (the one produced and the one from gold standard) automatically. This file gets read by `GoldStandardParser` and data are held as instances of `MedDocument`.

After the patient case files are read, assigned codes and found their relevant matches from the book, we compare the relevant matches with those from the gold standard. This gets done `EvaluationService` in the very end of the program run. Evaluation results are then stored in an XML file.⁵

3.5 Parsing the book

FIXME Martin

⁴/data/goldStandard.xml

⁵/data/output/evaluation.xml

- 4 Limitations given in the assignment. Present theories used. Are any theories outside of the curriculum been used? Explain the reason for the selection
- 5 How have the theories been adapted to the assignment?
- 6 Summary of the results.
- 7 Evaluation and discussion of the results. Limitations of your evaluation results.
- 8 Thoughts of potential improvements.
- 9 Conclusions.
- 10 Future work

Our system is accurate but in the future some refinements could be done to improve our query engine that we didn't manage to implement due to time constraints. For example an evaluation function that gives more importance to the hierarchy of the classes, for example if a code is subClassOf another code (E10, subclass and E10-14, the superclass in ICD for example). Since we use *Jena* to read the OWL ontologies, could be done through a simple graph traversal using suitable methods offered by the APIs, or using ad hoc SPARQL queries.

We could even taken better into account the synonyms: if two codes have the same synonyms, the may refer to the same disease.